

Skincancer HAM-dataset using Pytorch

Standardimporter

```
[1] import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader, Dataset
from torchvision import datasets, transforms, models
from torchvision.utils import make_grid
from torch.utils.data import WeightedRandomSampler

import os
from PIL import Image
# from skimage import io, transform

import numpy as np
import pandas as pd

from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import confusion_matrix,
classification_report
from sklearn.model_selection import train_test_split

import matplotlib.pyplot as plt
%matplotlib inline
```

```
[2] # Nedan används om man t.ex. vill ha tusentalsavgränsare:
import locale
locale.setlocale(locale.LC_ALL, '')
```

```
'LC_CTYPE=en_US.UTF-8;LC_NUMERIC=sv_SE.UTF-8;LC_TIME=sv_SE.UTF-8;LC_COLLATE=en_US.UTF-8;LC_MONETARY=sv_SE.UTF-8;LC_MESSAGES=en_US.UTF-8;LC_PAPER=sv_SE.UTF-8;LC_NAME=sv_SE.UTF-8;LC_ADDRESS=sv_SE.UTF-8;LC_TELEPHONE=sv_SE.UTF-8;LC_MEASUREMENT=sv_SE.UTF-8;LC_IDENTIFICATION=sv_SE.UTF-8'
```

Hjälpfunktioner

```
[3] def create_filename(filename):  
    global file_name  
    file_name = filename + "_e" + str(epochs) + "_bsz" +  
    str(batchsz) + \  
        "_lr" + str(f'{learning_rate:.0e}')    return file_name
```

```
[4] # https://pytorch.org/docs/master/notes/serialization.html  
  
def save_trained_model(modelname):  
    model_folder = "trained_models"  
    model_file_suffix = ".pt"  
    create_filename(modelname) # spottar ur sig ett filnamn i  
    variabeln "file_name"  
  
    full_model_filename = model_folder + "/" + file_name +  
    model_file_suffix  
  
    torch.save(model.state_dict(), full_model_filename)
```

```
[5] def load_trained_model(modelname):  
    model_folder = "trained_models"  
    model_file_suffix = ".pt"  
    create_filename(modelname) # spottar ur sig ett filnamn i  
    variabeln "file_name"  
  
    full_model_filename = model_folder + "/" + file_name +  
    model_file_suffix  
  
    if torch.cuda.is_available():  
        model.load_state_dict(torch.load(full_model_filename))  
    else:  
        model.load_state_dict(torch.load(full_model_filename,  
        map_location=torch.device('cpu')))
```



```
[6] # Följande återställer modellens vikter  
# mellan körningar:  
  
# usage: model.apply(weights_init)  
  
def weights_init(m):  
    if isinstance(m, nn.Conv2d):  
        torch.nn.init.xavier_uniform_(m.weight.data)
```

Importera data och definiera sökvägar

```
[7] imageFolder = "../../../ml/Datasets/skin-cancer-mnist-ham10000/images_in_one"

metadataSkincancerFilename = "../../../ml/Datasets/skin-cancer-mnist-ham10000/csv/HAM10000_metadata.csv"
```

```
[8] # Importera metadatan i en Pandas DataFrame:

skincancer_df = pd.read_csv(metadataSkincancerFilename)
```

Utforska och bearbeta vårt data

```
[9] # Kolla om vi har några noll-värden i vår Dataframe:

skincancer_df.isnull().sum()
```

```
lesion_id      0
image_id       0
dx             0
dx_type        0
age           57
sex            0
localization   0
dtype: int64
```

```
[10] # Enligt ovan är det bara i "age"-kolumnen som vi har noll-
värden.
# Dessa fyller vi ut genom att beräkna medelvärdet:

skincancer_df['age'].fillna((skincancer_df['age'].mean()),
inplace=True)
```

```
[11] # Definiera var vi har våra labels:

labels = skincancer_df['dx']
```

```
num_classes = len(labels.unique()) # --> 7
```

```
[12] # Innan vi fortsätter måste vi göra om
      # text-datan till numeriska features.

      # 1. Flytta labels sist i df
      # 2. Flytta age till efter image_id
      # 3. Alla categorical features emellan.

      skincancer_df = skincancer_df[['lesion_id', 'image_id', 'age',
                                      'dx_type', 'sex', 'localization', 'dx']]
```

```
[13] dxtype_feat_df = skincancer_df.iloc[:, 3]
      dxtype_feat_df.value_counts()
```

```
histo      5340
follow_up   3704
consensus    902
confocal     69
Name: dx_type, dtype: int64
```

```
[14] sex_feat_df = skincancer_df.iloc[:, 4]
      sex_feat_df.value_counts()
```

```
male      5406
female    4552
unknown    57
Name: sex, dtype: int64
```

```
[15] loc_feat_df = skincancer_df.iloc[:, 5]
      loc_feat_df.value_counts()
```

```
back      2192
lower extremity  2077
trunk      1404
upper extremity  1118
abdomen    1022
face       745
chest      407
foot       319
unknown    234
neck       168
scalp      128
hand       90
ear        56
```


```
genital          48
acral             7
Name: localization, dtype: int64
```

```
[16] label_df = skincancer_df.iloc[:, 6]
      label_df.value_counts()
```

```
nv      6705
mel     1113
bkl     1099
bcc      514
akiec   327
vasc     142
df       115
Name: dx, dtype: int64
```


```
[17] le = LabelEncoder()
      dxtype_feat = le.fit_transform(dxtype_feat_df)
      plt.hist(dxtype_feat)
```

```
(array([ 69.,  0.,  0., 902.,  0.,  0., 3704.,  0.,  0.,
        5340.]),
 array([0. , 0.3, 0.6, 0.9, 1.2, 1.5, 1.8, 2.1, 2.4, 2.7, 3. ]),
 <a list of 10 Patch objects>)
```




```
[18] loc_feat = le.fit_transform(loc_feat_df)
      plt.hist(loc_feat)
```

```
(array([1029., 2192., 463., 745., 319., 138., 2077., 296., 1404.,
        1352.]),
 array([ 0. ,  1.4,  2.8,  4.2,  5.6,  7. ,  8.4,  9.8, 11.2, 12.6, 14.
        ]),
 <a list of 10 Patch objects>)
```




```
[19] sex_feat = le.fit_transform(sex_feat_df)
     plt.hist(sex_feat)
```

```
(array([4552.,    0.,    0.,    0.,    0., 5406.,    0.,    0.,    0.,
        57.]),
 array([0. , 0.2, 0.4, 0.6, 0.8, 1. , 1.2, 1.4, 1.6, 1.8, 2. ]),
 <a list of 10 Patch objects>)
```



```
[20] labels = le.fit_transform(label_df)
     plt.hist(labels)
```

```
(array([ 327.,  514.,    0., 1099.,    0.,  115., 1113.,    0., 6705.,
        142.]),
 array([0. , 0.6, 1.2, 1.8, 2.4, 3. , 3.6, 4.2, 4.8, 5.4, 6. ]),
 <a list of 10 Patch objects>)
```



```
[21] le_skincancer_df = skincancer_df.copy()
      del skincancer_df
```

```
[22] le_skincancer_df['dxtype'] = dxtype_feat
      le_skincancer_df['sex'] = sex_feat
      le_skincancer_df['loc'] = loc_feat
      le_skincancer_df['label'] = labels
      le_skincancer_df.drop(columns=['localization', 'dx_type', 'dx'],
                             inplace = True)
      le_skincancer_df.head()
```

	lesion_id	image_id	age	sex	dxtype	loc	label
0	HAM_0000118	ISIC_0027419	80.0	1	3	11	2
1	HAM_0000118	ISIC_0025030	80.0	1	3	11	2
2	HAM_0002730	ISIC_0026769	80.0	1	3	11	2
3	HAM_0002730	ISIC_0025661	80.0	1	3	11	2
4	HAM_0001466	ISIC_0031633	75.0	1	3	4	2

```
[23] # Vi normaliserar metadata-features
      # innan vi skickar in det till Pytorch
      # för träning:

      from sklearn.preprocessing import MinMaxScaler

      scaler = MinMaxScaler()
```

```
[24] le_skincancer_df[['age', 'dxtype', 'loc']] =
      scaler.fit_transform(le_skincancer_df[['age', 'dxtype', 'loc']])
```

```
[25] le_skincancer_df.drop(columns='lesion_id', axis = 1,  
    inplace=True)
```

```
[26] le_skincancer_df.head(n = 5)
```

	image_id	age	sex	dxtype	loc	label
0	ISIC_0027419	0.941176	1	1.0	0.785714	2
1	ISIC_0025030	0.941176	1	1.0	0.785714	2
2	ISIC_0026769	0.941176	1	1.0	0.785714	2
3	ISIC_0025661	0.941176	1	1.0	0.785714	2
4	ISIC_0031633	0.882353	1	1.0	0.285714	2

```
[27] features_df = le_skincancer_df.iloc[:, :5]
```

```
[28] labels_df = le_skincancer_df.iloc[:, 5:]
```

Bygg en egen dataloader-klass för bilder och Pandas-dataframe

```
[29] class SkinCancerHamDF(Dataset):  
  
    def __init__(self, features_df, labels_df, root_dir,  
        transform = None):  
  
        self.features_df = features_df  
        self.labels_df = labels_df  
        self.root_dir = root_dir  
        self.transform = transform  
  
    def __len__(self):  
        return len(self.features_df)  
  
    def __getitem__(self, idx):  
        if torch.is_tensor(idx):  
            idx = idx.tolist()
```



```

img_name = self.features_df.iloc[idx, 0]
#print(type(img_name))
img_name = str(img_name)
#print(type(img_name))

full_img_name_woext = os.path.join(self.root_dir, \
                                    img_name)

full_img_name = full_img_name_woext + ".jpg"

# pillow:
image = Image.open(full_img_name)

# labels:
label = self.labels_df.iloc[idx]
label_np = np.array([label], dtype = int)

# features:
metadata = self.features_df.iloc[idx, 1:5]
metadata_np = np.array([metadata], dtype = float)

# Skapar en batch utan transforms:
sample = {'image': image, \
          'metadata': metadata_np, \
          'label': label_np}

# Skapar en batch med transforms:
if self.transform:
    sample = {'image': self.transform(image), \
              'metadata': metadata_np, \
              'label': label_np.flatten()}

return sample

```

```

[30] # Definiera en batch-storlek:
      batchsz = 2*5

      # Definiera vilken augmentation som ska göras:
      train_data_transform = transforms.Compose([

          transforms.Resize([224, 224]),\
                                                    transforms.ToTensor(),

          transforms.Normalize(mean = [0.485, 0.456, 0.406],

                                std = [0.229, 0.224, 0.225])

                                                    ])

      val_data_transform = transforms.Compose([

```

```

transforms.Resize([224, 224]),\
                                transforms.ToTensor(),
\
transforms.Normalize(mean = [0.485, 0.456, 0.406],
std = [0.229, 0.224, 0.225])
])

```

Dela upp i tränings- och valideringsdata

```

[31] features_train, features_val, labels_train, labels_val =
train_test_split(features_df, \

                labels_df, \

                test_size = 0.1, \

                shuffle = True)

```

```

[32] ham10k_train = SkinCancerHamDF(features_df = features_train, \
                                   labels_df = labels_train, \
                                   root_dir = imageFolder, \
                                   transform = train_data_transform)

```

```

[33] ham10k_val = SkinCancerHamDF(features_df = features_val, \
                                  labels_df = labels_val, \
                                  root_dir = imageFolder, \
                                  transform = val_data_transform)

```

```

[34] ham10k_train_dl = DataLoader(ham10k_train, \
                                 batch_size = batchsz, \
                                 num_workers = 4, \
                                 pin_memory = True, \
                                 )

```

```

[35] ham10k_val_dl = DataLoader(ham10k_val, \
                                batch_size = batchsz, \
                                pin_memory = True, \
                                num_workers = 4, \
                                )

```

```
[36] # Flytta till grafikkortet:

use_cuda = torch.cuda.is_available() # True/False
device = torch.device("cuda:0" if use_cuda else "cpu")
print(device)

# Override:
# device = "cpu"
```

```
cuda:0
```

Importera en Alexnet-modell

```
[37] # W_out = (W_in - Kernel_Filtersz + 2*padding) / stride + 1
# Avrundas neråt.

metadata_input_dim = features_df.shape[1] - 1 # första kolumnen är
bildnamnet...
```

```
[38] model = models.AlexNet(num_classes = num_classes)
```

```
[39] print(model)
```

```
AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2,
2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2,
2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
```

```

(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceiling_mode=False)
)
(avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
(classifier): Sequential(
  (0): Dropout(p=0.5, inplace=False)
  (1): Linear(in_features=9216, out_features=4096, bias=True)
  (2): ReLU(inplace=True)
  (3): Dropout(p=0.5, inplace=False)
  (4): Linear(in_features=4096, out_features=4096, bias=True)
  (5): ReLU(inplace=True)
  (6): Linear(in_features=4096, out_features=7, bias=True)
)
)

```

```

[40] # Hämta en batch för att skicka igenom
      # ett nätverk:

      for x in ham10k_train_dl:
          x = x['image']
          print("Ursprunglig storlek: ", x.shape)
          break

```

Ursprunglig storlek: torch.Size([32, 3, 224, 224])

```

[41] x_0 = x

```

```

[42] x_1 = model.features(x_0)

```

```

[43] x_2 = model.avgpool(x_1)

```

```

[44] x_2.shape

```

torch.Size([32, 256, 6, 6])

```

[45] model = models.AlexNet(7)

```

```

[46] model.modules

```

```

<bound method Module.modules of AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2,
2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2,
2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=9216, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=4096, out_features=7, bias=True)
  )
)>

```

```

[47] class MultInput(nn.Module):
    def __init__(self):
        super(MultInput, self).__init__()

        #model = models.AlexNet(num_classes = 7)

        self.alexnet = models.AlexNet(num_classes = 7)
        #self.

        self.fcc = nn.Sequential(
            nn.Linear(6*6*256 + metadata_input_dim, 512),
            nn.Linear(512, 256),
            nn.Linear(256, 64),
            nn.Dropout(p = 0.3),
            nn.Linear(64, num_classes),

```

```

    )

    def forward(self, x1, x2):
        x1 = self.alexnet.features(x1)
        x1 = self.alexnet.avgpool(x1)

        x1 = x1.view(x1.size(0), -1)
        x2 = x2.view(x2.size(0), -1)

        x = torch.cat((x1, x2), dim = 1)
        x = self.fcc(x)

    return x

```

```

[48] # Definiera en instans av CNN():
mult_model = MultInput()

# Flytta till rätt device för träning:
mult_model.to(device)

```

```

MultInput(
  (alexnet): AlexNet(
    (features): Sequential(
      (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
      (1): ReLU(inplace=True)
      (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
      (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (4): ReLU(inplace=True)
      (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
      (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (7): ReLU(inplace=True)
      (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (9): ReLU(inplace=True)
      (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (11): ReLU(inplace=True)
      (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
    (classifier): Sequential(
      (0): Dropout(p=0.5, inplace=False)
      (1): Linear(in_features=9216, out_features=4096, bias=True)
      (2): ReLU(inplace=True)
      (3): Dropout(p=0.5, inplace=False)
      (4): Linear(in_features=4096, out_features=4096, bias=True)
      (5): ReLU(inplace=True)
    )
  )
)

```

```

        (6): Linear(in_features=4096, out_features=7, bias=True)
    )
)
(fcc): Sequential(
  (0): Linear(in_features=9220, out_features=512, bias=True)
  (1): Linear(in_features=512, out_features=256, bias=True)
  (2): Linear(in_features=256, out_features=64, bias=True)
  (3): Dropout(p=0.3, inplace=False)
  (4): Linear(in_features=64, out_features=7, bias=True)
)
)

```

```

[49] # Dubbelkolla att modellen flyttats till
     # grafikkortet eller ej (True/False):

```

```

next(mult_model.parameters()).is_cuda

```

True

```

[50] # Hur många parametrar har modellen
     # att träna?

```

```

trainableparameters = []
for param in mult_model.parameters():
    # trainableparameters = param.numel()
    trainableparameters.append(param.numel())

tot_params = np.sum(trainableparameters)

print(f'Antalet träningsbara parametrar är {tot_params:n} st')

```

Antalet träningsbara parametrar är 61 901 902 st

Återställ nätets vikter

```

[51] mult_model.apply(weights_init)

```

```

MultiInput(
  (alexnet): AlexNet(
    (features): Sequential(
      (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=
(2, 2))
      (1): ReLU(inplace=True)
      (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,

```

```

ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=
(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
)
(avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
(classifier): Sequential(
  (0): Dropout(p=0.5, inplace=False)
  (1): Linear(in_features=9216, out_features=4096, bias=True)
  (2): ReLU(inplace=True)
  (3): Dropout(p=0.5, inplace=False)
  (4): Linear(in_features=4096, out_features=4096, bias=True)
  (5): ReLU(inplace=True)
  (6): Linear(in_features=4096, out_features=7, bias=True)
)
)
(fcc): Sequential(
  (0): Linear(in_features=9220, out_features=512, bias=True)
  (1): Linear(in_features=512, out_features=256, bias=True)
  (2): Linear(in_features=256, out_features=64, bias=True)
  (3): Dropout(p=0.3, inplace=False)
  (4): Linear(in_features=64, out_features=7, bias=True)
)
)

```

Loss och optimiser

```

[52] # Definiera loss-function och vilken optimerare som ska användas:

epochs = 10 # verkar räcka med tanke på overfitting.

learning_rate = 1e-4

criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(mult_model.parameters(), lr =
learning_rate)

```



```
[53] modelname = create_filename("multi_model_cnn")
     print(modelname)
```

```
multi_model_cnn_e20_bsz32_lr1e-04
```

Trainingloop:

```
[54] mult_model.apply(weights_init)

do_validation = 1

import time
start_training_time = time.time()

num_training_images = len(ham10k_train_dl.dataset)
num_val_images = len(ham10k_val_dl.dataset)

accuracy_train = []
losses_train = []
accuracy_val = []
losses_val = []

for epoch in range(epochs):
    start_epoch_time = time.time()

    num_correct_per_epoch_train = 0
    tot_loss_per_epoch_train = 0.0

    for idx_train, inputs_train in enumerate(ham10k_train_dl):
        images_train = inputs_train['image'].float().to(device)
        metadata_train =
inputs_train['metadata'].float().to(device)
        labels_train = inputs_train['label'].flatten().to(device)

        mult_model.train()
        # Låt modellen göra förutsägelser:
        predictions_train = mult_model.forward(images_train,
metadata_train)

        # Beräkna statistik från träningen på träningsmängden:
        _, predicted_train = torch.max(predictions_train.data, 1)
        num_correct_per_batch_train =
torch.sum(predicted_train.detach() == labels_train.detach())
        num_correct_per_epoch_train +=
num_correct_per_batch_train
```

```

        # Beräkna loss:
        loss_per_batch_train = criterion(predictions_train,
labels_train)
        tot_loss_per_epoch_train += loss_per_batch_train.item() *
batchsz

        # Nolla ackadumulerade gradienter, göra
        # backprop. & uppdatera vår optimeringsfunktion:
        optimizer.zero_grad()
        loss_per_batch_train.backward()
        optimizer.step()

# Validation
if do_validation == True:
    start_eval_time = time.time()
    mult_model.eval()

    num_correct_per_epoch_val = 0
    tot_loss_per_epoch_val = 0.0

    for idx_val, inputs_val in enumerate(ham10k_val_dl):
        images_val = inputs_val['image'].float().to(device)
        metadata_val =
inputs_val['metadata'].float().to(device)
        labels_val = inputs_val['label'].flatten().to(device)

        with torch.no_grad():

            predictions_val = mult_model.forward(images_val,
metadata_val)

            _, predicted_val = torch.max(predictions_val, 1)
            # Behöver ej använda .detach med torch.no_grad()
            num_correct_per_batch_val =
torch.sum(predicted_val == labels_val)
            #num_correct_per_batch_val =
torch.sum(predicted_val.detach() == labels_val.detach())
            num_correct_per_epoch_val +=
num_correct_per_batch_val

            loss_per_batch_val = criterion(predictions_val,
labels_val)

            tot_loss_per_epoch_val +=
loss_per_batch_val.item() * batchsz

    end_eval_time = time.time()
else:
    pass

# Beräkna statistik från epoken:
# Use torch.Tensor.item() to get a Python number
# from a tensor containing a single value:

```

```

        # acc, train:
        accuracy_per_epoch_train = num_correct_per_epoch_train.item()
    / num_training_images
    accuracy_train.append(accuracy_per_epoch_train)

    # loss, train:
    loss_per_epoch_train = tot_loss_per_epoch_train /
num_training_images
    losses_train.append(loss_per_epoch_train)

    if do_validation == True:
        # acc, val:
        accuracy_per_epoch_val = num_correct_per_epoch_val.item()
    / num_val_images
        accuracy_val.append(accuracy_per_epoch_val)

        # loss, val:
        loss_per_epoch_val = tot_loss_per_epoch_val /
num_val_images
        losses_val.append(loss_per_epoch_val)
    else:
        pass

    end_epoch_time = time.time()
    epoch_time = end_epoch_time - start_epoch_time

    # epoch startar på 0, därav "+1" nedan:
    if do_validation == True:
        print(f"Epok {epoch+1:03}, {epoch_time:4.1f} sek. ---
train acc = {accuracy_per_epoch_train:4.3f} --- val acc =
{accuracy_per_epoch_val:4.3f} --- train loss =
{loss_per_epoch_train:4.5f} --- val loss =
{loss_per_epoch_val:4.5f}")
    else:
        print(f"Epok {epoch+1:03}, {epoch_time:4.1f} sek. ---
train acc = {accuracy_per_epoch_train:4.3f} --- train loss =
{loss_per_epoch_train:4.5f}")

    end_training_time = time.time()

    delta = end_training_time - start_training_time

    print(f'\nTraining took {delta/60:.2f} minutes.')

```

```

Epok 001, 38.7 sek. --- train acc = 0.672 --- val acc = 0.680 --- train
loss = 0.95033 --- val loss = 0.91445
Epok 002, 37.4 sek. --- train acc = 0.693 --- val acc = 0.703 --- train
loss = 0.84901 --- val loss = 0.85084
Epok 003, 38.5 sek. --- train acc = 0.710 --- val acc = 0.714 --- train
loss = 0.79567 --- val loss = 0.81534
Epok 004, 38.3 sek. --- train acc = 0.731 --- val acc = 0.728 --- train

```

```
loss = 0.74376 --- val loss = 0.76473
Epok 005, 39.1 sek. --- train acc = 0.750 --- val acc = 0.752 --- train
loss = 0.69561 --- val loss = 0.71512
Epok 006, 38.2 sek. --- train acc = 0.764 --- val acc = 0.754 --- train
loss = 0.65553 --- val loss = 0.69711
Epok 007, 38.1 sek. --- train acc = 0.775 --- val acc = 0.756 --- train
loss = 0.61099 --- val loss = 0.70397
Epok 008, 39.0 sek. --- train acc = 0.793 --- val acc = 0.763 --- train
loss = 0.57082 --- val loss = 0.71992
Epok 009, 36.8 sek. --- train acc = 0.812 --- val acc = 0.750 --- train
loss = 0.51579 --- val loss = 0.77799
Epok 010, 36.5 sek. --- train acc = 0.830 --- val acc = 0.746 --- train
loss = 0.46456 --- val loss = 0.82143
Epok 011, 37.8 sek. --- train acc = 0.848 --- val acc = 0.749 --- train
loss = 0.41798 --- val loss = 0.78295
Epok 012, 36.6 sek. --- train acc = 0.864 --- val acc = 0.745 --- train
loss = 0.37439 --- val loss = 0.89099
Epok 013, 37.4 sek. --- train acc = 0.881 --- val acc = 0.754 --- train
loss = 0.34085 --- val loss = 0.88456
Epok 014, 36.3 sek. --- train acc = 0.887 --- val acc = 0.745 --- train
loss = 0.32284 --- val loss = 1.00964
Epok 015, 37.2 sek. --- train acc = 0.908 --- val acc = 0.749 --- train
loss = 0.25535 --- val loss = 1.11338
Epok 016, 36.6 sek. --- train acc = 0.917 --- val acc = 0.742 --- train
loss = 0.23726 --- val loss = 1.01239
Epok 017, 37.9 sek. --- train acc = 0.930 --- val acc = 0.723 --- train
loss = 0.19908 --- val loss = 1.10928
Epok 018, 36.8 sek. --- train acc = 0.935 --- val acc = 0.741 --- train
loss = 0.18606 --- val loss = 1.11537
Epok 019, 36.8 sek. --- train acc = 0.937 --- val acc = 0.735 --- train
loss = 0.17785 --- val loss = 1.38238
Epok 020, 38.2 sek. --- train acc = 0.946 --- val acc = 0.729 --- train
loss = 0.15765 --- val loss = 1.19798
```

Training took 12.54 minutes.

```
[55] save_trained_model(modelname)
```

Visualisering av träningsomgången


```
[56] if do_validation == False:

    plt.figure(figsize = (12, 7))
    plt.title("Visualisering träningsomgången")
    plt.plot(losses_train, label = "Loss")
```

```

plt.plot(accuracy_train, label = "Andel korrekta")
plt.legend()
plt.show()
else:
    plt.figure(figsize = (12, 7))
    plt.title("Visualisering träningsomgången")
    plt.plot(losses_train, label = "Loss, train")
    plt.plot(accuracy_train, label = "Acc train")
    plt.plot(losses_val, label = "Loss, validation",
linestyle='dashed')
    plt.plot(accuracy_val, label = "Acc validation",
linestyle='dashed')
    plt.legend(loc = 'lower left')
    plt.show()

```



Utvärdering

```

[57] def evaluate_model(num_eval_images, data_generator, model):

    model.eval()

    start_eval_test_time = time.time()

    global labels_val_true_np, predicted_val_np

    labels_ground_truth = []
    predicted_classes_val = []

```

```

with torch.no_grad():
    correct_val = 0
    # Antal iterationer = num_images / batchsz = x st.
    for inputs in data_generator:
        images = inputs['image'].float().to(device)
        metadata = inputs['metadata'].float().to(device)
        labels = inputs['label'].flatten().to(device)

        # Spara för mer utvärdering utanför denna funktion:
        labels_ground_truth.append(labels.cpu().numpy())

        # Låt modellen göra förutsägelser:
        predictions_val = mult_model.forward(images,
metadata)
        _, predicted_val = torch.max(input = predictions_val,
dim = 1)

        predicted_classes_val.append(predicted_val.cpu().numpy())

        correct_val += (predicted_val == labels).sum()

    # Platta till listorna...
    labels_ground_truth = [item for sublist in
labels_ground_truth for item in sublist]
    predicted_classes_val = [item for sublist in
predicted_classes_val for item in sublist]

    #flat_list = []
    #for sublist in predicted_classes_val:
    #    for item in sublist:
    #        flat_list.append(item)

    labels_val_true_np = np.array(labels_ground_truth)
    predicted_val_np = np.array(predicted_classes_val)

    end_eval_test_time = time.time()
    eval_test_time = end_eval_test_time -
start_eval_test_time

    print(f'Test accuracy: {correct_val.item()}/{num_eval_images}
= {100*correct_val.item()/(num_eval_images):5.2f} %')
    print(f"\nEvaluation took {eval_test_time:.2f} seconds.")

```

```

[58] evaluate_model(len(ham10k_val_dl.dataset), ham10k_val_dl,
mult_model)

```

Test accuracy: 730/1002 = 72.85 %

Evaluation took 4.15 seconds.

```
[59] cm_pytorch = confusion_matrix(labels_val_true_np,
    predicted_val_np)
```

```
[60] print(cm_pytorch)
```

```
[[ 9  2 15  0  0  5  0]
 [10 20 16  3  2  9  0]
 [ 8  2 66  1 14 21  0]
 [ 1  0  6  4  2  2  0]
 [ 4  5 20  0 66 29  0]
 [ 1  2 49  1 35 556  0]
 [ 0  2  0  0  2  3  9]]
```

```
[61] print(classification_report(labels_val_true_np,
    predicted_val_np))
```

	precision	recall	f1-score	support
0	0.27	0.29	0.28	31
1	0.61	0.33	0.43	60
2	0.38	0.59	0.46	112
3	0.44	0.27	0.33	15
4	0.55	0.53	0.54	124
5	0.89	0.86	0.88	644
6	1.00	0.56	0.72	16
accuracy			0.73	1002
macro avg	0.59	0.49	0.52	1002
weighted avg	0.75	0.73	0.73	1002

```
[ ]
```

```
[ ]
```